

Context-Aware-VR

Mathias Ooms

University of Antwerp

september 2020

context-extraction

Context characteristics:

- ▶ Eye tracking
- ▶ Position tracking

Why reinforcement learning

Supervised learning

- ▶ needs data to increase performance
- ▶ can never become better than the quality of the data

Why reinforcement learning

Reinforcement learning

- ▶ provides a way to have an agent learn completely by itself (without requiring data)
- ▶ agents can become better than the 'best' players

Reinforcement learning

Policy Network

- ▶ transforms the input to the possible actions

Policy Gradients

- ▶ method to train a policy network

Policy Gradients

- ▶ starts off from a completely random network




Figure 1: The principle of RL

- ▶ calculate for each possible action a probability
- ▶ agent will choose an action and get feedback (reward)
- ▶ the goal of the agent is to receive as much reward as possible

Policy Gradients

How does it improve the Policy network?

- ▶ In the first generations it will miserably fail, until the agent gets lucky (received a good reward)
- ▶ this successful series of actions now need to be more preferred over other ones
- ▶ when receiving a good reward we multiply the probability with a positive gradient
- ▶ when receiving a bad reward we multiply the probability with the same negative gradient

"Troubles in paradise"

Unfortunately there are also downsides:

- ▶ *Credit assignment problem*: punished hard for small mistakes
- ▶ Not appropriate in complex environments